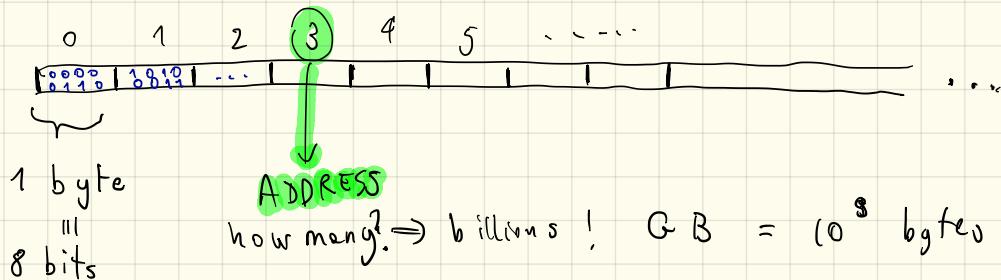


# STORING ARRAYS IN MEMORY

The computer needs to write somewhere all the info that needs to be processed/stored

RAM → linear container



convention,  
all the  
computers  
we nowadays  
organized this  
way

All the bits always have a state 1 / 0, depending on past computations, random initialization, ...

What is the most natural way of storing a sequence of numbers in this memory?

E.G.: we want to store

[ 2, 19, 3, 8 ]

## Binary representation

2 → 0 0 0 0 0 0 1 0

14 → 0 0 0 0 1 1 1 0

3 → 0 0 0 0 0 0 1 1

8 → 0 0 0 0 1 0 0 0



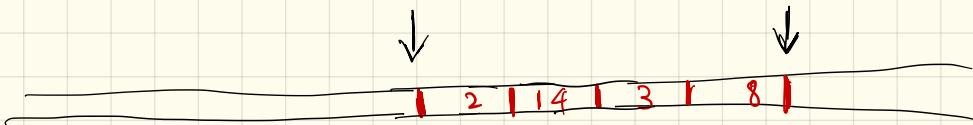
they all fit  
in 1 byte!

for example

take 14 →

8 bits							
2 <sup>7</sup>	2 <sup>6</sup>	2 <sup>5</sup>	2 <sup>4</sup>	2 <sup>3</sup>	2 <sup>2</sup>	2 <sup>1</sup>	2 <sup>0</sup>
128	64	32	16	8	4	2	1
0	0	0	0	1	1	1	0

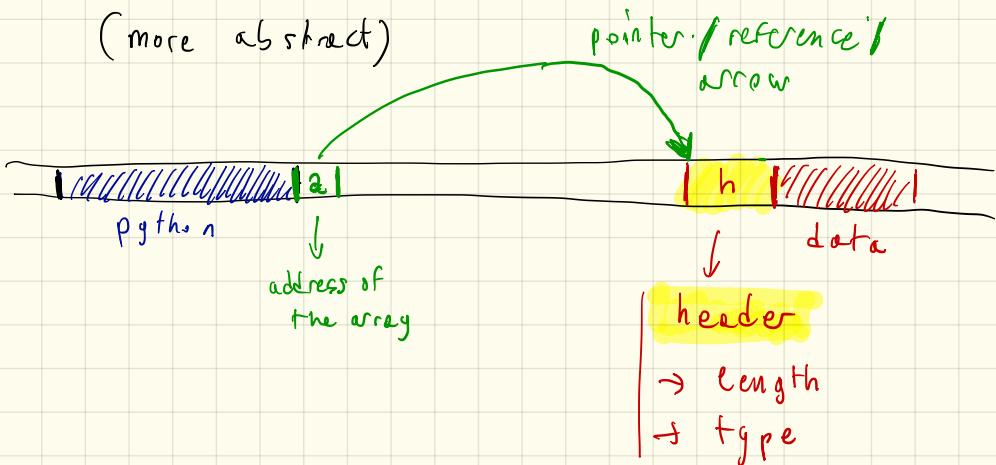
⇒ Use consecutive slots in the RAM :



start? end?

there is nothing special about these bits,  
unless I have a flag or something that  
points there!

# Storing 1-d ndarrays (#1)



WHY DO I NEED THE TYPE?

- `uint8` ⇒ each element takes 1 byte
- `uint32` ⇒ each element takes 4 bytes
- ⇒ homogeneity of the entries +  
length +  
type =  
= now the CPU knows how to read the content!

# Storing python List (#1)

Could I use the same compact representation for lists?

header	data
1 h	2   14   (3)   8   ...

flexible → swap 3 for  $\underbrace{3}_{100}$

ASSIGNMENT

The operation would  
be  $O(N)$ , where  
 $N$  is the length of the  
vector

→ I need more  
bits to store

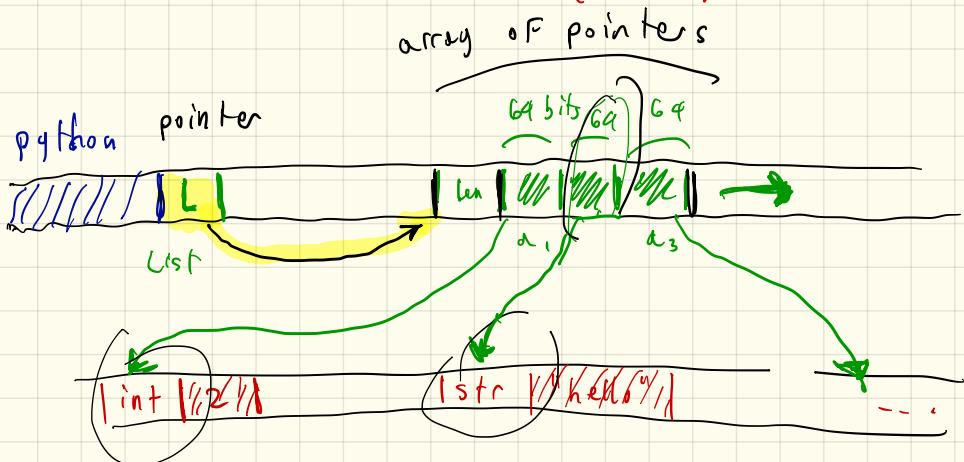
it → I have  
to move all  
the rest of  
the entries

→ we need  $O(1)$   
⇒ constant cost!

flexible  
TYPE → what if I swap 3 for  
"a", how does the program  
know how to interpret the bits?

IT DOES NOT WORK

## Storing python list (#2)



NOW I NEED 2 JUMPS TO FETCH THE  
CONTENT OF THE LIST  $L$ , BUT  
assignment becomes  $O(1)$

- write new content somewhere else
- update the pointer

why is it efficient to do operations within same type?

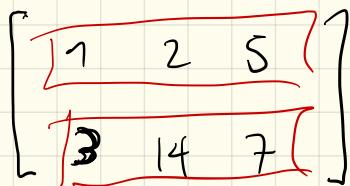
$$\underbrace{\text{int64}}_{\text{same operation?}} \quad 1 + 2.5 = 3.5$$

$$\underbrace{\text{float64}}_{\text{NO}} \quad 1.0 + 2.5 = 3.5$$

First row requires  
a conversion ( $\text{int} \rightarrow \text{float}$ )

same result

# Storing 2-d ndarrays (#1)



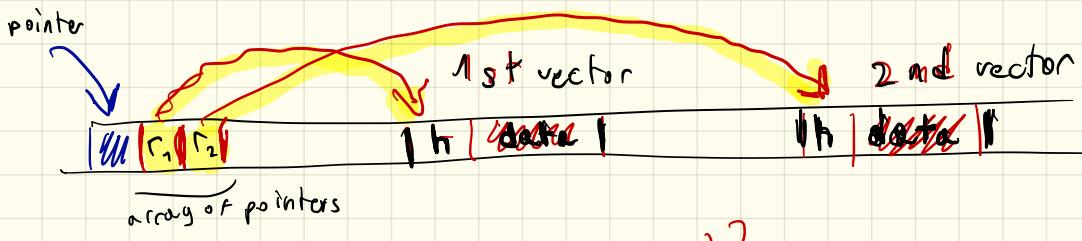
2D object :

How to share  
it?

(RAM is 1D)

(NOT HOW NUMPY DOES IT!)

1st possible way  $\rightarrow$  store each row as a separate vector, and then link them



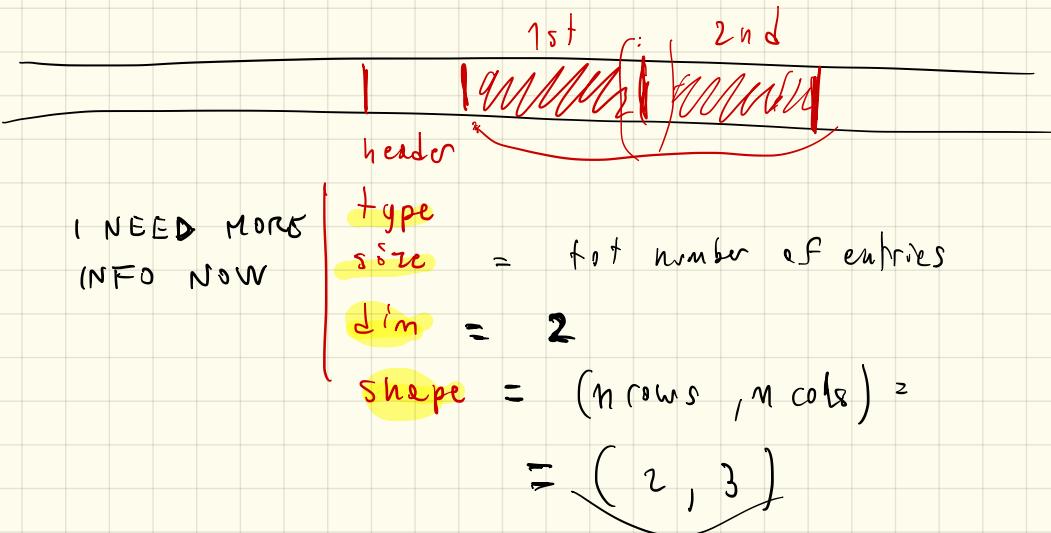
$$[[1\ 2\ 5], [3\ 4\ 7]]$$

this is indeed similar to how Python would store a list of lists

add another layer of  
pointers

# Storing 2-d ndarrays (#1)

Numpy groups all the data in contiguous slots  $\Rightarrow$  SINGLE HEADER



In what order am I storing the table?

Row-major ORDER  $\rightarrow$  C, C++, python, ...  
Column-major ORDER  $\rightarrow$  Fortran, Matlab, Julia, ...

Row-major  $\Rightarrow$  go row by row!